


OF@TIEN+ IoTCloud

Toward Large-Scale Microservice Orchestration with Kubernetes
Technology in IoTcloudServe@TEIN Testbed Infrastructure

Krerk Piromsopa, Ph.D.
Computer Engineering
Chulalongkorn University





Facebook Page

<https://www.facebook.com/iotcloudserve/>



Github

<https://github.com/IoTcloudServe>

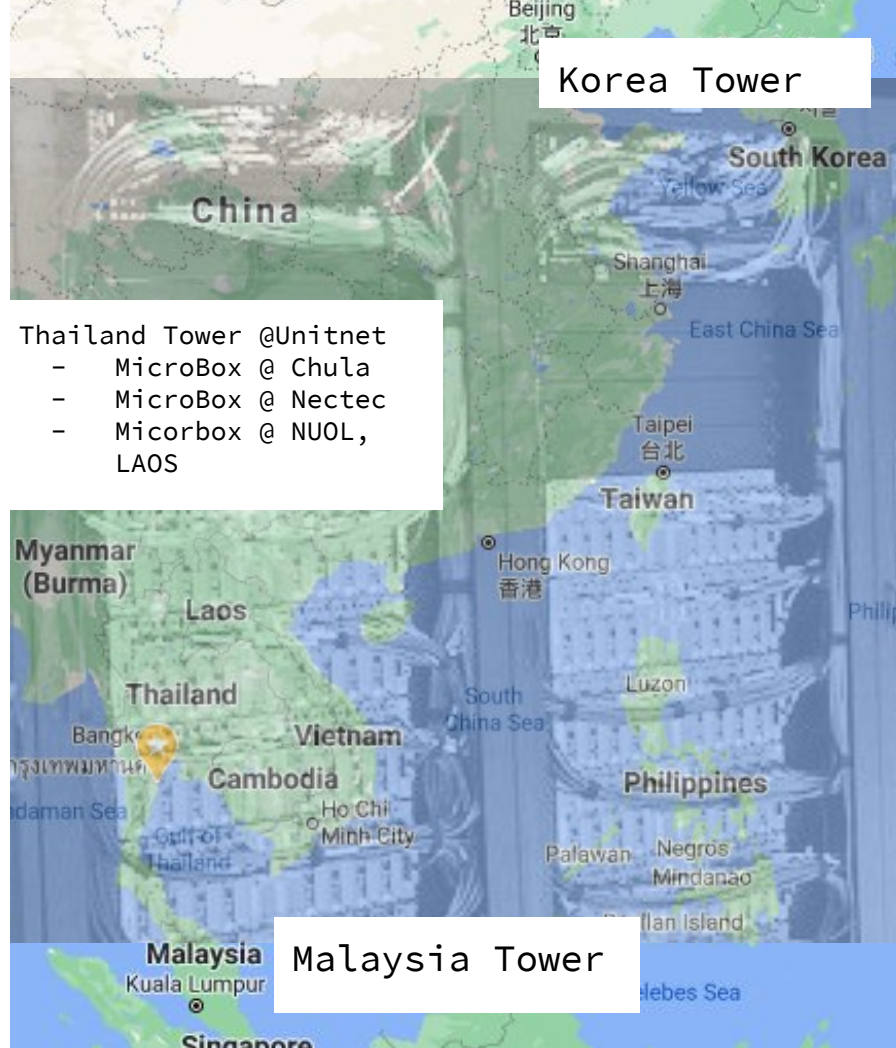
IoTCloud

— — —

- ★ Infrastructure & Applications
- ★ Large-Scale Microservice Orchestration
(Kubernetes Technology)
- ★ Research & Directions
 - Monitoring/Logging/Visibility
 - Control/Policy (Quota)
 - Storage Management
 - Federation Concepts
- ★ Summary

Infrastructure

- ★ Tower & Storage
- ★ Microboxes for developments and edge application
- ★ ...



Several Opportunities for Research and Developments

— — —

★ Management

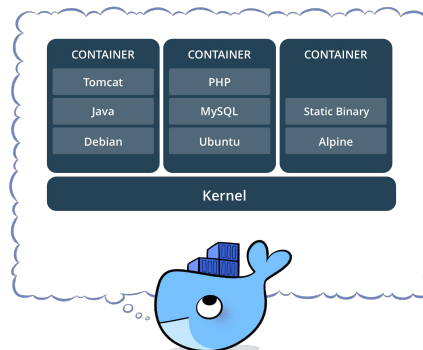
- Microservices Orchestration
- Quota/Control Policy
- Storage Management
- Federation
- Edge & Cloud

★ Applications

- Smart Agriculture/Farm
- Smart Energy
- Smart Mobility
- ... Smart X

Choices of Software/System Management

- ★ Application - Docker/Container
- ★ Orchestration - Kubernetes
- ★ Interface/Control - Rancher
- ★ Storage - R00K ceph
- ★ Monitor - Grafana + Prometheus
- ★ Logging - EFK
- ★ Service Mesh - ISTIO
- ★ Edge/Small Cluster - K3S



K3S



Istio



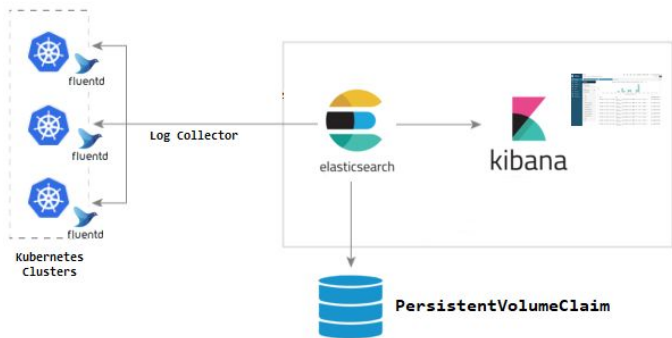
More details in the following session.

Monitor/Logging/Visualization

— — —

Logs

- ★ Collector - Fluentd
- ★ Storage - Elasticsearch
- ★ Visualization - Kibana

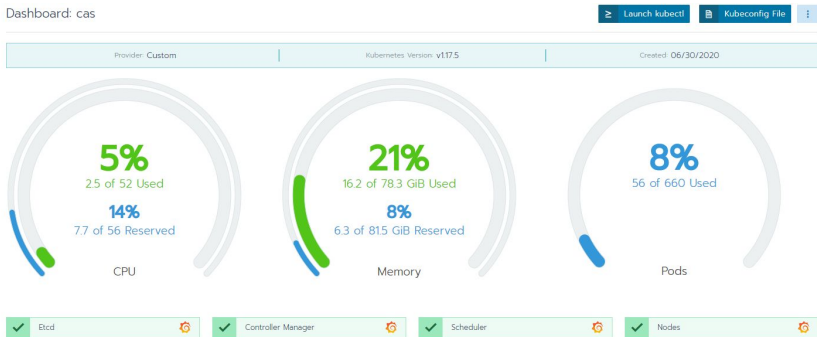


Metrics

- ★ Collector - Prometheus
- ★ Visualization - Grafana
- ★ (built-in with rancher)



Dashboard: cas



Control/Policy

— — —

More details in the following session.

▼ Resource Quotas
Configure how much of the resources the project can consume

Resource Type	Project Limit	Namespace Default Limit
CPU Limit ▼	32000	32000
CPU Reservation ▼	e.g. 2000	e.g. 500
Memory Limit ▼	e.g. 2048	e.g. 1024
Memory Reservation ▼	e.g. 2048	e.g. 1024
Storage Reservation ▼	e.g. 50	e.g. 10
Services Load Balancers ▼	e.g. 50	e.g. 10
Services Node Ports ▼	e.g. 50	e.g. 10
Pods ▼	e.g. 50	e.g. 10
Services ▼	e.g. 50	e.g. 10

+ Add Quota

- ★ Standard Kubernetes cluster comes name spaces and quota (via admission plugins)
- ★ Rancher introduces projects and quotas for both project and namespaces.
- ★ A namespace is assigned to a project.

Storage Management

- ★ Ceph supports Shared Files (CephFS), Block (RBD), and Objects (S3).
- ★ Others
 - Longhorn (block only)
 - Minio (Object only)
 - NFS (file only) - not self manage



Microservice Orchestration

Microservice Orchestration in a cluster

- ★ Ingress (nginx), mapping different paths to different services
- ★ Use Horizontal Pod Autoscaler for each service.
- ★ Use Database and/or CephFS for sharing.

This is practical for small number of services in a cluster.

Microservice Orchestration in a cluster (ctd.)

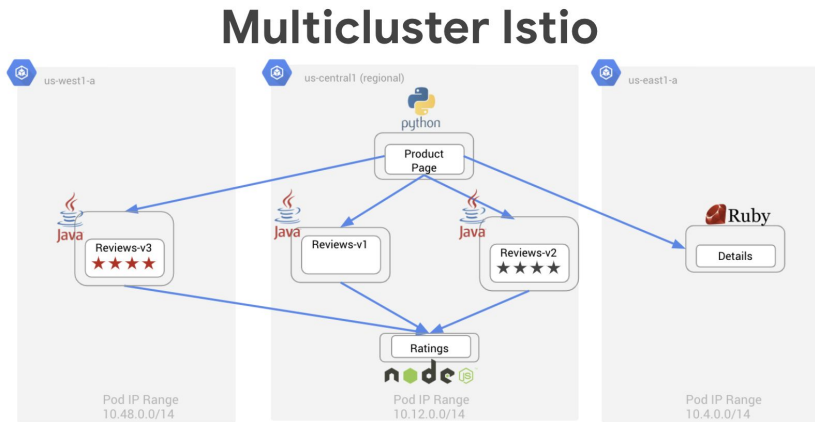
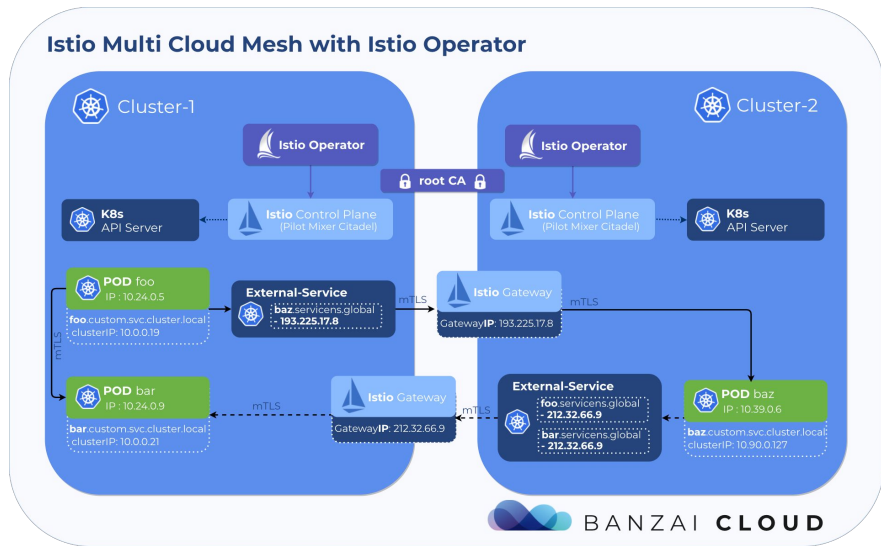
- ★ Mapping path - share cookies/sessions
- ★ Each service can scale independently.
- ★ Limited to a cluster.

```
s.co/api/account > account-service  
s.co/api/auth > auth-service  
s.co/api/exam > exam-service  
s.co/api/file > file-service  
s.co/ > frontend  
s.co/api/exim > exim-service  
s.co/api/mail > mail-service  
s.co/api/job > job-service  
s.co/api/form > form-service
```

More details in the following session.

Federation

- ★ Each kubernetes cluster is operated separately.
- ★ Use ISTIO gateway to mesh cross-cluster applications.



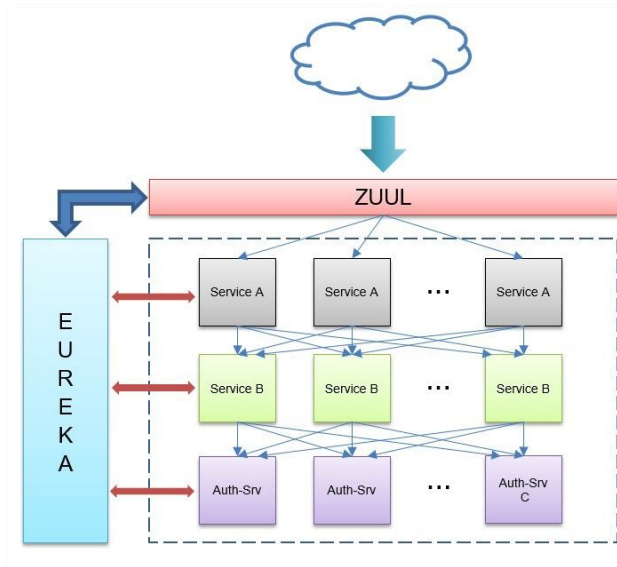
Federation (ctd.)

- ★ Istio mesh solves the issue at network layer.
- ★ Application can still be accessed from the same URL without (major) software modification.

Alternative solution(s)

- ★ Netflix ZUUL + Eureka
- ★ Dynamic scale regardless of cluster/cloud providers
- ★ Require software modification

- ★ Apigee (from Google), Kong (opensource), etc...



**We are now capable
of managing
kubernetes clusters.**

Best Practices for Cluster Management

- ★ HA Rancher installation with 2 x K3S nodes + 1 mysql
 - Use haproxy as a load balancer for K3S (rancher)
- ★ Provision K8S from rancher
- ★ Or Add existing cluster to rancher



K3S



Best Practices for Cluster Management (ctd.)

★ Databases

- MariaDB (Galera master-master) + haproxy
- MongoDB (replica set + sharding)

★ Storage

- Rook CEPH
(CephFS for sharing, RBD for high performance, Object for S3)

★ Scaling

- Small Applications - ingress + HPA
- Large Applications - ingress + HPA + ISTIO
- Multi-cluster Applications - ingress + HPA (+ ZUUL + eureka)

**This is just the
beginning.**

Ideas

— — —

- ★ Cloud-native storage with support for edge/off-line services (eg. IoT application, ad-hoc services)
- ★ Managing edge node with kubernetes (KubeEdge?)
- ★ more

Thank you
Q&A